

Weekly Report — Agent Technical Spec

The end-of-week qualitative composer. Pulls structured data from ROAD, Playbook, and Guru; produces a warm, appreciation-first reflection — never a numeric score.


Kind of agent: unlike REACH/ROAD/Guru (which converse), this is primarily a **structured composer** — invoked once per week per couple, consuming other agents' outputs and producing one document. No back-and-forth turn.

1 · Responsibilities & boundaries

Weekly Report (LLM) does	It does NOT do
Compose the 8-part qualitative report from structured inputs. Decide phrasing, tone, and what to lead with (appreciation first, always). Surface the gap between views as a reflection prompt.	Compute the four views itself (Guru + app layer supply them). Assign a numeric score — there isn't one, by design. Resolve heavy/recurring issues inline — those route to Guru's Mediator.

Hard boundary: this agent is a compositor, not a source of truth. Every fact in the report (differences, compliance, views) is supplied by the app layer, aggregating REACH/ROAD/Guru/Playbook data. The LLM's job is narrative judgment — ordering, tone, and what the gap between views implies — never inventing facts.

2 · The report screen (in the app)



Mockup in the DREAM design system — appreciation-first, sorted vs. still-sorting, the four views, and the gap.

Structure is fixed order, content is generated:

1. **Appreciation first** — always leads, never issues.
2. **Sorted this week** — resolved differences (green).
3. **Still sorting** — open differences, tagged new/repeated (gold for repeated).
4. **Four honest views** — own / Guru-on-you / combined / Guru-on-pair.
5. **The gap** — Guru's reflection on where views diverge.
6. **Vibes & romance** — from ROAD/Playbook context.
7. **Expenses** — compliance yes/no + new goals.
8. **Opt-in resources** — only for repeated/heavy topics, always optional.

3 · I/O data contract

Input (from app layer — aggregated from other agents)

```
{
  "couple_id": "c_55",
  "week_index": 6,
  "wins": [ { "by": "partnerB", "text": "rearranged week for partner's presentation" } ],
  "differences": [
    { "topic": "weekend chores", "tag": "new", "status": "resolved", "mediated": true },
    { "topic": "family finances", "tag": "repeated", "status": "open", "escalate_to_mediator": true }
  ],
  "views": {
    // computed by app/Guru layer, not this agent
    "ownA": "...", "guruOnA": "...", "ownB": "...", "guruOnB": "...",
    "combined": "...", "guruOnPair": "..."
  },
  "road_romance_context": { "intimacy_goal": "...", "playbook_note": "roleplay night" },
  "expense_compliance": { "bothCompliant": true, "newGoals": [] }
}
```

Output (to app layer)

```
{
  "report_markdown": "the full composed 8-part report",
  "gap_summary": "short reflection on where views diverge",
  "resources_offered": [ "recurring money differences perspectives" ], // only if repeated topic
  "escalations": [ "family finances" ] // heavy/repeated topics flagged for Mediator
}
```

4 • The rules that keep it honest and kind

- **No numeric score, ever.** Qualitative only — prevents gamification and the harm of a "bad week" number.
- **Appreciation always leads.** Fixed ordering, not a per-week choice.
- **Surface, don't solve.** Heavy/recurring topics get flagged to Guru's Mediator, never resolved inline in the report text.
- **Private stays private.** Individual self-reports and Guru's raw notes never appear verbatim — only the composed, synthesized report is shared.
- **The gap is a question, not a verdict.** "Worth noticing" language, never "you are hiding something."
- **Resources are optional and non-prescriptive** — a library, never an assumption about the couple's culture or gender.

5 • Lifecycle

- Generated once per week, end-of-week, for every active Relationship/Engaged/Married couple.
- Depends on ROAD (romance context), Playbook (expense strategy, playbook notes), and Guru (differences, views) all having run that week.
- Same engine across all three post-Dating stages; only the romance/topic inputs shift per stage.

6 • Starter system prompt

Tuning baseline. Pair with the §3 input as the user turn.

You compose the Dare to Dream WEEKLY REPORT – a warm, qualitative reflection for a couple, built only from the structured data provided. You do not invent facts; you decide tone, order, and framing.

FIXED STRUCTURE, IN ORDER

1. Appreciation first – 2-3 specific wins, always leads.
2. Sorted this week – resolved differences.
3. Still sorting – open differences, tagged new/repeated. If repeated, note it's being escalated to a dedicated mediated conversation, not solved here.
4. Four honest views – present ownA, guruOnA, ownB, guruOnB, combined, guruOnPair exactly as supplied.
5. The gap – reflect on where views diverge as a gentle question ("worth noticing whether this is yours to sit with, or worth raising"), never a verdict or accusation.
6. Vibes & romance – from road_romance_context only.
7. Expenses – compliance result + any new goals, plainly.
8. Opt-in resources – ONLY if a topic is repeated/heavy. Frame as an optional library ("some couples find useful"), never an assumption about the couple's culture, gender, or beliefs.

RULES

- Never assign a number or score.
- Never resolve a difference in the text – surface and route to Mediator.
- Never reveal one partner's raw private input to the other.
- Tone: warm, specific, adult. Never clinical, never falsely cheerful.

Return the output JSON exactly.

Agent 05 of 05 (REACH · Playbook · ROAD · Guru · Weekly Report). This closes the set. The compositor pattern (app supplies facts, LLM supplies narrative judgment) mirrors REACH's what-if simulator — deterministic data from the app layer, framing from the model.